

Week 2 Optimisation for Machine Learning

Neimhin Robinson Gunning, 16321701

February 13, 2024

1 Part (a)

1.1 Part (a) (i)

To represent the expression x^4 using sympy we first create a symbol object for x (line 1 in Listing 1), and then raise it to the fourth power using standard python arithmetic notation. The `**` operation is overloaded such that `x**4` yields an object representing the expression x^4 . Letting $y(x) = x^4$ the derivative dy/dx is found with `(x**4).diff()`. Running the code we find this derivative to be $4x^3$.

Listing 1: Source code to find the derivative of x^4 using sympy. The result is $4x^3$.

```
src/ai.py      Thu Jan 11 09:02:38 2024      1
1: import sympy as sp
2: x = sp.symbols('x')
3: y = x**4
4: print(y.diff()) # 4*x**3
5:
```

1.2 Part (a) (ii)

The analytic derivative of x^4 is compared, for a range of values of x , to the finite difference method approximations with a range of perturbations δ in Figure 1. The analytic derivative and the finite difference approximation with $\delta = 0.01$ are different. The function used to compute the finite difference approximation is presented in Listing 2. It takes a sympy function `f`, a specific x value `xval`, and a perturbation to compute the finite difference according to Equation 1.

$$\frac{\hat{dy}}{dx}(x) = \frac{(x + \delta)^4 - x^4}{\delta} \quad (1)$$

1.3 Part (a) (iii)

The plot in Figure 1 shows the range $0.001 \leq x \leq 0.0011$, and we shows that the error between the finite difference approximation and the analytic solution increases as the perturbation δ increases. Perturbations in $[0.001, 0.1]$ are shown¹. The relationship between the error and δ is not linear, and furthermore the relationship depends on the value of x , though this dependence is not visible at the scale in the plot.

¹Including perturbations in $[0.001, 1]$ would make the plot less readable so a narrower range was chosen

Listing 2: Ellided source code to estimate the derivative of x^4 using the finite difference method. The function `diff_with_pert` can be used to estimate the derivative of an arbitrary function of x , such that $f(x + \delta)$ and $f(x)$ are defined. The full script is in the appendix and named `src/a11.py`.

```
src/finite_diff_eg.py      Wed Jan 31 14:28:44 2024      1
1: import sympy as sp
2: # finite difference
3: def diff_with_pert(f, xval, pert=0.01):
4:     global x
5:     delta_x = pert
6:     return (f.subs(x, xval + delta_x) - f.subs(x, xval)) / (delta_x)
7: x = sp.symbols('x')
8: y = x**4
9: dydx = y.diff()
10: analytic_ys = [dydx.subs(x, i) for i in my_range()]
11: # ...
12: for pert in np.array([0.01, 0.1, 0.15]):
13:     dydx_finite = diff_with_pert(y, x, pert=pert)
14:     # ...
15:
```

The accuracy of the finite difference approximation of the derivative is greater for smaller values of δ . The reason for this is related to the definition of the derivative. The derivative is defined as

$$\lim_{\delta \rightarrow 0} \frac{f(x + \delta) - f(x)}{x + \delta - x}$$

, i.e. the derivative is the function approached by this expression as δ approaches 0. As we decrease δ and apply the finite difference method we are approaching the true derivative.

2 (b)

2.1 (b) (i)

Listing 3 shows ellided source code for a python class that can be used to run gradient descent. Ellided functions allow for specification of the function to be optimized, the convergence condition, the max number of iterations etc. Line 13 calculates the gradient at the current value x_{n-1} , and line 14 updates the value based on the gradient and step size, $x_n := x_{n-1} - dy/dx(x_{n-1})\alpha$. The convergence condition is a function of x_n and x_{n-1} . Examples of this class in use are in the appendix, e.g. in `src/b1.py`, in which the gradient function is computed with `sympy`.

2.2 (b) (ii)

In Figure 2 is a plot of gradient descent on x^4 with $x_0 = 1$ and $\alpha = 0.1$. We see that updates to $|\hat{x}|$ and $y(\hat{x})$ are initially large but quickly become tiny. Because x^4 is very flat around the minimum (see subplot (c)) we have a very small gradient, and thus very small step sizes, so the algorithm updates extremely slowly in the flat region of x^4 , i.e. near $\hat{x} = 0$, but cannot reach it $\hat{x} = 0$ in a finite number of iterations.

2.3 (b) (iii)

In Figure 3 several runs of gradient descent on x^4 are visualized. Subplot (a) shows $|\hat{x}|$ for each iteration, which is the distance of the estimate \hat{x} from the optimal value, $x = 0$. With initial value $x = 1$ and step size $\alpha = 0.1$ we




Figure 1: A comparison of the analytic derivation of $dy/dx = 4x^3$ and some approximations of dy/dx using the finite difference method.

Listing 3: Ellided source code a python class that implements gradient descent with a fixed step size.

```
src/gradient_descent_listing.py    Wed Jan 31 15:38:46 2024    1
1: class GradientDescent():
2:     # ...
3:     def iterate(self):
4:         import math
5:         x_value = self._start
6:         old_x_value = None
7:         iteration = 0
8:         while True:
9:             yield [iteration, float(x_value), float(self._function(x_value))]
10:            iteration += 1
11:            if self._max_iter > 0 and iteration > self._max_iter:
12:                break
13:            grad_value = self._gradient(x_value)
14:            x_value -= self._step_size * grad_value # Update step
15:            if old_x_value is not None and self._converged(x_value, old_x_value):
16:                yield [iteration, float(x_value), float(self._function(old_x_value))]
17:                print("converged")
18:                break
19:            old_x_value = x_value
```

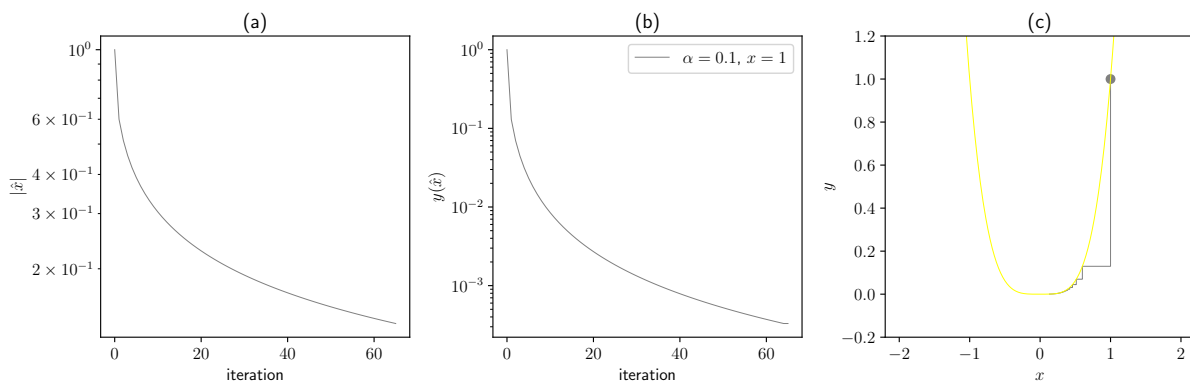


Figure 2: Running gradient descent on x^4 with initial value $x_0 = 1$ and step size $\alpha = 0.1$. The yellow line in subplot (c) is x^4 . Note the log scales in (a) and (b).

Table 1: Various runs of gradient descent on x^4 with different outcomes. The maximum number of iterations was set to 100, and the convergence threshold to 0.001. Some of the α values were purposefully selected to lie near the margin of non-convergence, which can be calculated with $\alpha_m = \frac{2x_0}{4x_0^3}$. So rows 8 and 9 use $\alpha_m \pm \epsilon$.

	α	start	convergence time	final guess
0	0.01	1.0	100	0.3323366923926806
1	0.02	1.0	100	0.24119415115009837
2	0.03	1.0	97	0.20146846304616808
3	0.04	1.0	89	0.18252923619780276
4	0.05	1.0	83	0.1691984331910481
5	0.1	1.0	66	0.13379629115760108
6	0.4999999999999999	1.0	55	-0.07827643923095963
7	0.5	1.0	100	-1.0
8	0.7812500000000009	0.8	29	inf
9	0.7812499999999989	0.8	24	0.02474194758833489
10	0.8	0.7	33	-0.06609857197723801

see that the steps are initially large, i.e. moving quickly to the optimum, but on successive iterations the steps slow down dramatically. In subplot (b) we see the value of $y(\hat{x})$ on each iteration.

The convergence criterion for gradient descent adopted here is when $|\alpha\Delta x| < 0.001$ or x is $\pm\infty^2$. For the function x^4 we can see from Table 1 that the final guess, i.e. the estimate of $\arg \min x^4$ is most accurate and converges fastest with larger values of α , **unless** α is too large and the iterations of gradient descent degenerate, i.e. subsequent iterations give worse or unchanged estimates. The reason for this is that the gradient of x^4 is tiny as x approaches the minimum. So, a smaller value of α means the convergence condition, $|\alpha\Delta x| < 0.001$, can be met with a higher value of Δx , but in the case of x^4 , Δx is higher when x is further from the optimal value 0. In other words a lower α allows stopping further from the minimum. This is why the final guess is worse for smaller values of α , but why does a smaller α lead to slower convergence time?

Since α is constant for a single run of gradient descent, the convergence condition depends on Δx , namely we need $|\Delta x| < 0.001/\alpha$. A larger α entails a larger step, but there are two factors at play here:

1. A larger step size may entail a smaller Δx on the subsequent iteration.
2. A larger step size moves us away from meeting the convergence condition.

It turns out for x^4 that the first of these factors is more important, in cases where gradient descent is converging. Thus larger α results in faster convergence (for x^4).

Similarly, the initial value of x turns out to be relatively less important, assuming the configuration converges. The difference in progression between starting with any of $x = 0.7, 1, 1.5$, while holding $\alpha = 0.1$, is small, as seen in Figure 3. The algorithm needs a handful of extra iterations when x_0 is larger, but arrives at the same result and follows a similar trajectory.

As we increase α there are two important points at which behaviour changes. The first point is the perfect setting of α such that the minimum is found after one iteration. Above this point the iterations of gradient descent will chatter, but still converge, up until the second important point, where the first iteration results in a worse estimate. After this point the algorithm does not converge. (In the case of x^4).

²Meaning the floating point values `math.inf` or `-math.inf`

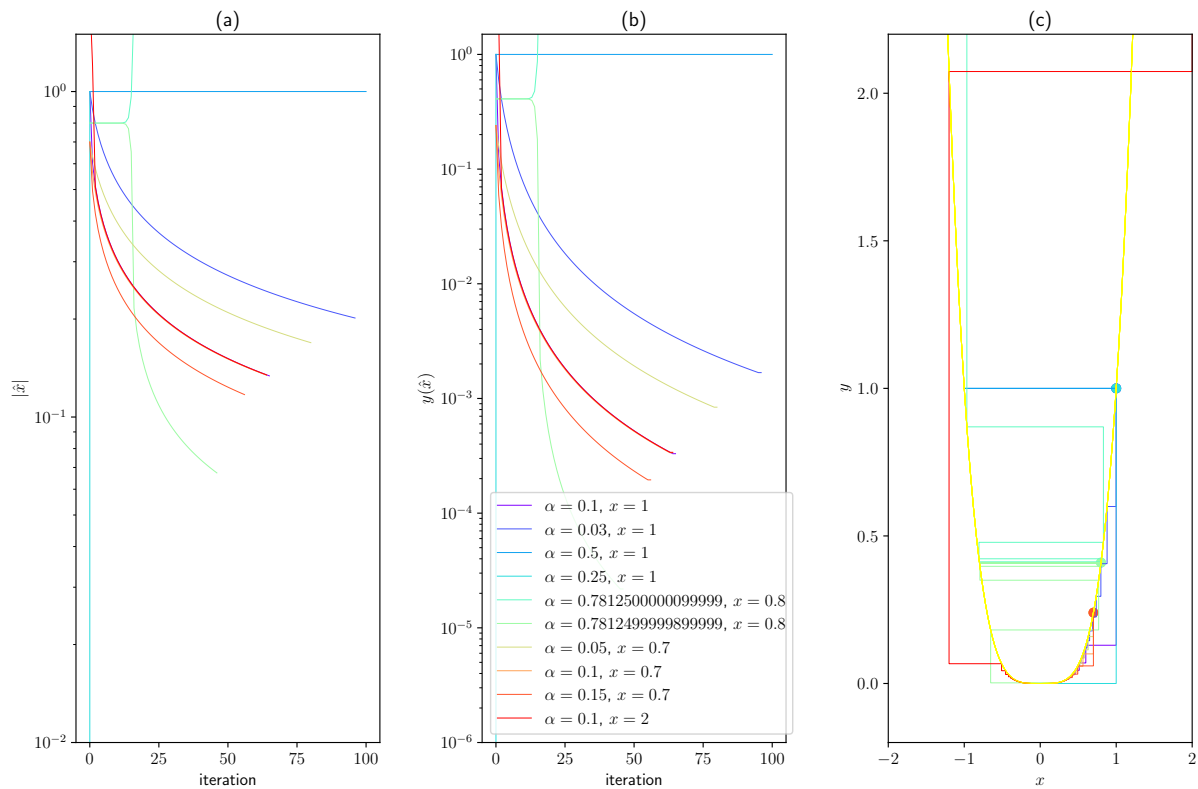


Figure 3: Various runs of gradient descent on x^4 . One run, where $\alpha = 0.5$ and $x_0 = 1$, is stuck in a 2-step loop, flipping \hat{x} back and forth between 1 and -1. The yellow line in subplot (c) is x^4 .

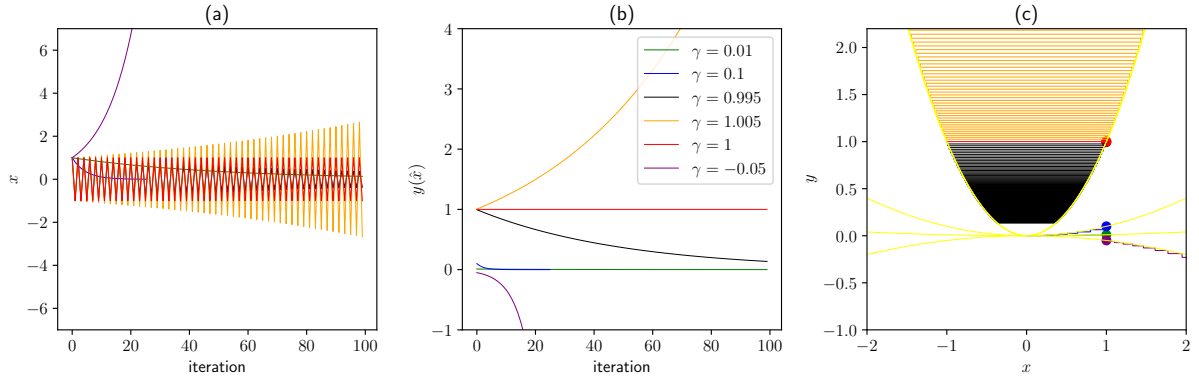


Figure 4: Various runs of gradient descent on γx^2 . Subplot (a) shows the estimate of x for each iteration. The yellow lines in subplot (c) are plots of the various γx^2 functions.

3 (c)

3.1 (c) (i) Optimising γx^2

Firstly, note that if γ is negative, then γx^2 and $\gamma|x|$ have no minimum, and thus gradient descent will behave weirdly. For instance, in my implementation, with enough iterations it will actually finish with `math.inf`, which is reasonable. If $\gamma = 0$ then gradient descent terminates immediately, and any value for x minimizes $\gamma x^2 = 0x^2$. Increasing γ increases the magnitude of the gradient, resulting in larger steps, but the convergence time and accuracy will then depend on the initial value.

Several runs of gradient descent on γx^2 are visualized in Figure 4.

Let $d_0 = |x_0 - \arg \min_x \gamma x^2|$ be the initial distance to the optimal value of x . Let's assume x_0 is 'far' from

$$\arg \min_x \gamma x^2$$

, i.e. d_0 is large, meaning gradient descent will at first take several steps in the same direction. Increasing γ to γ' , i.e. $\gamma' > \gamma$, results in a larger gradient $2\gamma'x'_0 > 2\gamma x_0$ meaning the first step is larger. If after the first iteration the distance has improved, $d_1 < d_0$ then gradient descent will converge, if not it will diverge.

Let's assume we are in a converging condition. While the second gradient $2\gamma'x'_1$ is not necessarily larger than $2\gamma x_1$, the updated x'_1 is now closer to the minimum than x_1 . The result of these dynamics is that with a larger γ' there is a fewer number iterations until the first 'overshoot'. This means that convergence time is smaller for the larger γ' than for γ . Also, for fixed x and α , the step size is greater when γ is greater which means we can move closer to the minimum $x = 0$ without triggering the convergence condition, $\alpha \frac{dy}{dx}(x) \leq \epsilon$. Thus, both the convergence time and the final estimate are better when γ is higher, assuming a convergence condition.

3.2 (c) (ii) Optimising $\gamma|x|$

Here we assume $\gamma > 0$. In Figure 5 several runs of gradient descent on $\gamma|x|$ are plotted. The magnitude of the gradient of $\gamma|x|$ is constant for all $x \neq 0$, $dy/dx = \gamma$. The function $\gamma|x|$ is not smooth because it is not differentiable at $x = 0$. Here we assume $\frac{d\gamma|x|}{dx}(0) = 0$. The only cases where we see convergence are when the algorithm lands exactly on $x = 0$, because the step size is constant, so when there's an overshoot the next iteration will revert back to the previous estimate. In all other cases there will be "chatter" once the estimate is close enough to the minimum, namely $|\hat{x}| < \alpha\gamma$. At this point the algorithm will flip between two estimates repeatedly until `max_iter`.

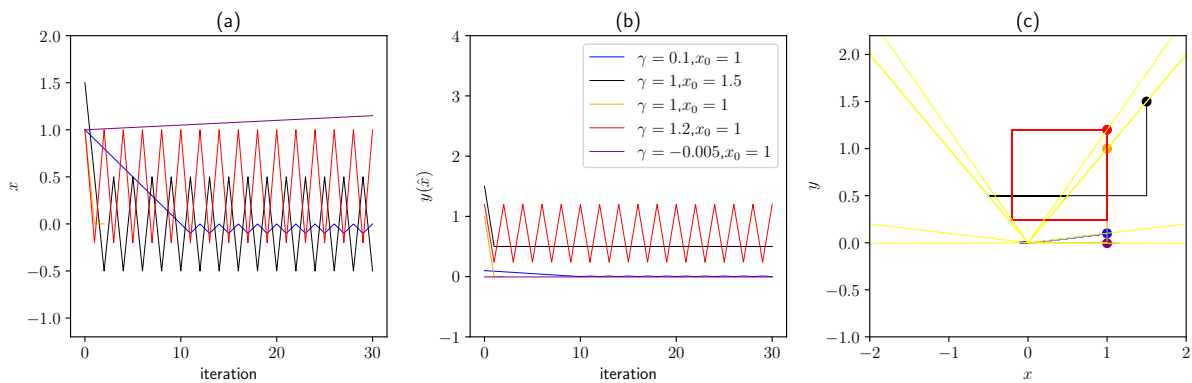


Figure 5: Various runs of gradient descent on $\gamma|x|$. Subplot (a) shows the estimate of x for each iteration. The yellow lines in subplot (c) are plots of the various $\gamma|x|$ functions.

Therefore, rather than discuss convergence time as defined previously we can consider the number of iterations until the algorithm loops. For fixed x_0 and α the time to first loop is negatively correlated with γ , because a larger γ means larger steps. The error of the final estimate (given sufficient iterations), d , will satisfy $d < \alpha\gamma$, thus a smaller γ entails a smaller maximum error. If we tweaked the algorithm to record and return the best choice of x for the whole run then the maximum error would be $\frac{1}{2}\alpha\gamma$.